

Word Vector Representations for Syntactically Disparate Event Descriptions

Travis Wolfe

Johns Hopkins University

travis@cs.jhu.edu

Abstract

In this work we present an unsupervised method for learning word vector representations that are conducive to identifying coreferent mentions of events. This is a hard problem because event mentions are weakly tied to syntax and can be realized in many ways. We discuss the justifications for a model that decomposes coherence into syntactic and conceptual aspects and provide a training procedure for it. An evaluation framework is provided but without results due to difficulty with numerical optimization.

1 Introduction

Knowledge representations for events in natural language has been given many different treatments. Much linguistic work does not concern itself with the problem of learning and the representation on which they study linguistic regularities. In NLP representations are often comprised of ad hoc feature and lists of trigger words which are only available as long as there is an expert able to write them. This data is often task specific and difficult to acquire. Some annotated corpora have been created with frames or propositions labelled in hopes that machine learning techniques might generalize this supervision to new data. In this work, we propose an unsupervised model for learning event representations based at the lexical level based on as little supervision and as much data as possible.

Our model of words decomposes words into the conceptual properties of the predicates and arguments they evoke/describe and syntactic properties

of the words themselves. As is generally true with unsupervised learning, you cannot learn things that are not present in the data, so this goal of learning event representations comes hand in hand with a hypothesis about the distribution of event descriptions and their arguments. The hypothesis is that there is regularity in the realization of predicates and their arguments which is independent of syntax which can be observed in local fragments of text. Previous work in learning features for words has operated in the paradigm of language modeling, and the features they learn are interesting but constrained (c.f. related work). The main drawback is that these representations do not allow generalization across syntactic category. This ability is paramount to modeling events because they can be described in a wide variety of syntactic realizations. Take Figure 1 for example, which gives three descriptions of a wedding, and note that 1) there is no clear syntactic correspondence between the descriptions and 2) a good amount of information is captured near the event's trigger word in each case.

The claim is not that arguments are always close, but that they are biased towards being close¹ and when they are not close, the intervening words are not systematically idiosyncratic to the event/predicate being described.

Thus we are beginning to describe a statistical theory of coherence based on locality of predicate and argument trigger words. What remains is to define mathematically how this coherence should be mod-

¹Grice's Maxim of Manner suggests that the speaker should be brief. If the speaker is to be both brief and informative, related concepts must appear close together.

| phrase | event trigger category | ARG0 category | ARGM category |
|--------------------------------------|------------------------|---------------|---------------|
| Jane’s joyous wedding | NN | NNP+POS | JJ |
| Jane was joyous when she got married | VB | NNP+PRP | JJ |
| Jane joyously gave her vows | VBD+PRP\$+VBZ | NNP | RB |

Figure 1: three syntactically disparate descriptions of the same event

elled.

2 Model

A primary goal of this work is to learn from unannotated text and the simplest way to do this is via language modeling, i.e. inducing a score or probability distribution over word sequences. Previous work has shown that we can learn very fine grained information by inducing features for a language model (c.f. related work). While previous work has referred to this as language modelling, this has the unwanted connotation that we are interested in the numerical scores that come out of the language model for some NLP task. Instead, we’ll refer to our modelling as addressing *coherence*, and we’ll discuss the factors that lead to coherence and how we might learn them. Later in section 5, we’ll discuss ways in which our model of coherence fails as a general explanation of coherence and how we might improve it.

Lets start with a simple model of coherence where we look at a small phrase of words and decide if it is coherent. Since we don’t know, or rather want to learn, how words collocate and compose, we will impose very few restrictions on how the words might interact in determining coherence. Lets use a single hidden layer neural network to describe this process of determining coherence². The inputs to the network will be one vector for every word in the phrase and the output will be a coherence score (which we will refer to as $s(\cdot)$). The hidden layer in the middle will store some notion of non-linear composition, but we will not assume we know how to parse.

This network has the nice property that we can propagate a gradient all the way from the output

score back to the values in the word vectors. We can then learn features for words by training the network to accomplish a discriminative task that looks like comprehension. Given a phrase x and a corrupted (and hence incoherent) phrase \tilde{x} , we want to train our model to recognize that the score of the uncorrupted phrase is more coherent than the corrupted one: $s(x) \geq s(\tilde{x}) + 1$. This training procedure, provided we can get coherent phrases and an automatic method of corruption, means that we can train on as much unlabelled training data as we have.

The model just described closely resembles the model of (0), and can learn fine grained (seemingly) semantic similarities like “Apple” is similar to “Dell” if you take similarity to mean cosine distance of the vectors learned for “Apple” and “Dell”. The problem with this method is that you will learn that “marry” and “wedding” are very dissimilar. The reason for this is that they are from different syntactic categories and you would never use one as a direct replacement for the other. If we want a more general notion of similarity, or in this case type-level event coreference³, then we need to relax the assumption of replacability in a sequence of words with a less syntax dependent notion of coherence.

One way to do this is to remove the order of the words in the sequence (this is used in the CBOW method of (0)), but this throws the baby out with the bathwater. With no order you cannot tell the difference between “dog bites man” and “man bites dog”, and surely you would learn a different conception of “biting” if you thought people went around doing it to dogs. A more general way to state this is that syntax plays a role in semantic interpretation, and we should not remove it entirely from our model of coherence.

Another way to remove the syntactic dependence

²Note that while describing the properties of this network we will refer to this as a black box computational unit that has capabilities to do things like compose and argmax vectors. While it is outside of the scope of this work to explain the mechanisms by which this occurs, single layer feed forward networks are universal function approximators (0) and can represent such computations with enough paramters and training data.

³type-level event coreference in the sense that instances may not be coreferent, the phrase “John married Mary” describes a different event from “John and Mary’s wedding”, but holding the all of the arguments constant, coreference holds.

is to consider that words that describe events that have different syntactic category live in different places of the vector space we are learning our representation in, but that there is some kind of mapping between event descriptions. For example we might want to learn a transformation on our vector space that takes the vector for “marry” (in the verbal part of the space) to a vector near “wedding” (in the nominal part of the space).

We tested the plausibility of a model like this by trying to learn this type of transformation on labelled pairs of nominal and verbal trigger words for events like “testament” and “testify”. This data was taken from NOMLEX (0), which provides us with 3,135 noun-verb pairs. We took the embeddings of (0) and tried to learn a linear map that minimized the Euclidean distance between the vector representation for the noun x_n and the projected vector for the verb x_v :

$$\min_P \sum_{(n,v) \in \text{NOMLEX}} \|x_n - Px_v\|_2^2$$

When we compared the error of this model (on training data), it was not substantially lower than the variance from the mean of the nominal vectors, $\sum \|x_n - E[x_n]\|_2^2$, which indicates that this model has no hope of relating this data in an accurate way. Nonlinear transforms were not attempted due to the size of the data and the apparent difficulty of the mapping.

On second inspection it doesn’t seem like this method of projection would work for a few reasons. First, since the word representations have not been preprocessed for polysemy, projection is much harder than it would be if we assumed that we were projecting a centroid in some conceptual space⁴. If a word vector needs to store two meanings, e.g. “testament” as in “old” or as in what you do in court, then the projection would need to preserve this polysemy, and in a way that represents the syntactic category information, even if that notion might be null in some sense (e.g. there is no verb that corresponds to the “old testament” sense of “testament” because it is not an event).

⁴this is assuming that the values in a word vector can represent either a point in space (a.k.a. a centroid) or a distribution like one that you might get out of a feature vector in a log-linear model

We then return to our simpler model of coherence, but try to force the vectors we learn for words to be less syntax dependent. We can do this by separately modeling the syntactic and conceptual aspects separately. Lets define a word vector v_i to be a composition of a vector representing its syntactic properties s_i and its non-syntactic properties w_i .

$$v_i = s_i \oplus w_i$$

As before, we will build up a phrase vector x by stacking the words in a phrase, $x = [v_{i-k}, \dots, v_i, \dots, v_{i+k}]$, for some small context size k . We want to come up with a way of learning s_i in a way that will favor syntax and be tied between various words (else we have something that is equivalent to the first model), so we will represent s_i as a row in a matrix $S_{t(i), \bullet}$, where $t(i)$ is the part of speech tag of word i . In principle $t(i)$ could be more fine grained than a part of speech tag, but it should be easily determined by surface forms of the text so that it can be used on large unlabelled datasets. We can also arrange all of the w_i as rows in a matrix W .

With regard to how \oplus should be implemented, we could use multiplication, additions, or some other operator, but we favor vector stacking because it implies that there are dimensions that account for syntactic versus semantic coherence, a distinction which we intuitively believe to exist.

3 Learning

To train our model, we sought a large dataset that has many mentions of events, so we opt for a corpus comprised of newswire text taken from the New York Times, a part of the Annotated Gigaword dataset (0), ranging from 2000 to 2004. This yields a dataset with 153,762 unique words and 456,746,432 phrases. Given this corpus, we first give every word that appears at least 100 times an index and label all other words as OOV⁵. Numbers are replaced by a unique token that only says how many digits are in the token⁶. We then break text up in windows of size 5 ($k = 2$) and shuffle all of the instances.

We need to choose the corruption model carefully because it affects the parameters we learn. We follow (0) in that we only corrupt the middle word

⁵we will learn one vector for OOV even though there are many OOV words

⁶e.g. “120” \rightarrow 3D and “1995” \rightarrow 4D

of the language model, but we still have to decide whether to corrupt the word, its part of speech tag, or both. In this work we do not consider perturbations of both because that leads to a very diffuse gradient which updates both the W and S matrices rather than just one. As a heuristic for when to perturb the word vs it’s tag, we use

$$p(\text{perturb word}) = \frac{\sqrt{|W|} \times H(p_W)}{\sqrt{|W|} \times H(p_W) + \sqrt{|S|} \times H(p_S)}$$

where $|W|$ is the number of parameters in the W matrix, $H(\cdot)$ is the entropy of a distribution, and p_W and p_S are the unigram distributions of words and part of speech tags respectively⁷.

The feedforward network was implemented in Theano (0), a python library which provides automatic differentiation and support for GPU computing. Training was run using a variant of batched stochastic gradient descent called AdaGrad (0)⁸. We set the dimension of both the syntactic and non-syntactic vectors to be 32 and the hidden layer to be 40 (on the small end of the spectrum with respect to previous work, but the dimension tends to correlate with training time).

4 Evaluation

In order to evaluate our model, there are a few perspectives we can take. The first question, is does our compositional model of coherence do any better than the baseline model of one vector per word?

From the learning curves, it appears that the baseline method outperforms the two part model for the first 150 epochs. It is likely that this is due to search error because the baseline is nested in the full model (functional subset). It is not clear if more training time would alleviate this problem.

We also want to know about the nature of the syntax-stripped vectors we learn, in W . In this work we have regarded the conceptual or semantic properties of a word simply as the residual of what syntax doesn’t explain. This notion is a bit unhelpful in

⁷You could consider methods that make a perturbation choice based on the phrase, but these methods are substantially slower at training time or would require a lot of offline computation. A careful implementation could do this efficiently, but that is not in the scope of this paper. The method used in this paper only requires a few bits of entropy and no memory accesses.

⁸Batches of 500 examples worked best in our experiments




Figure 2: Held out data learning curves for the baseline model (black) and the decomposition model (blue). This is only the first few hours of training due to a last minute bug fix.

light of the fact that there are properties of words that we commonly refer to “semantic” properties, and we would like to know if we are learning a representation that speaks to these properties. To explore one aspect of this, we consider a low-resource classification task where the labels are semantic properties that do not seem to be related to the surface forms of the words used to describe them, namely the event class of *State*, *Process*, and *Accomplishment*⁹. Given only 45 verbs across these classes, the leave-one-out accuracy of a classifier trained on the representations of these words achieved 73% accuracy, significantly greater than would be achieved by guessing the most common class (36%)¹⁰. This indicates that whatever is being learned, it is correlated with at least one aspect of what we consider to be relevant properties to the semantics of events.

5 Shortcomings

Though we have discussed our hypothesis of how predicate argument descriptions appear in text, and

⁹data taken from <http://www.sfu.ca/person/dearmond/322/322.event.class.htm>

¹⁰the vectors for this experiment were taken from the baseline model

how our model might learn from these patterns, it is worth dwelling on the shortcomings of this model.

The most obvious objection is that this model is very local in its view of composition. Long distance dependencies pose a problem in this model with respect to argument interaction. If we assumed that arguments were independent conditioned on the predicate, and that with some probability arguments appeared in the same phrase as their predicate, then we might hope to learn the selectional preferences. However, if there are interactions between the arguments, and we consistently see only some of the arguments in the phrase, then the model has no way to tell if the argument is conditionally likely or not given the limited information in the phrase.

Related to this objection is the objection that there is no parsing or modeling of binary constituent unification. The easiest and perhaps least satisfying response to this is that parsing and models of larger spans of text are outside of our computational ability right now. To put this in perspective, even with a model that was chosen for its computational efficiency like this one, training high quality vectors can take three (0) to eight (0) *weeks* of CPU time. This is not to detract from the value of these training routines, which should be weighed against the many-years of effort required to come up with them. But training on tree structure or wider windows will likely take advances in either neural net training techniques or computational power.

Another problem with this model is that there is no explicit attempt to model instances of data rather than types. That is, we learn representations for words, but never reason explicitly about the meaning of a word in context aside from the state of the hidden layer of the network. This means that ambiguity, be it introduced by polysemy or non-determinism in composition, is reasoned about in an ad hoc and flat way. Some work has addressed this problem (0; 0).

Lastly, this model lives in an ungrounded paradigm of learning. Whatever it learns, it learns from the New York Times, and can never observe any of the events described. This means that it has no real sense of time or duration, but rather only represents these concepts as far as they help explain the text it sees. It is an open question of how much real world knowledge can be learned only from text and whether this can be learned without a system-

atic prior or bias that is present only in humans. The experiment concerning event class at least indicates that some things that are not obvious, like telicity, show up in a systematic enough way that the model can learn something correlated with it.

6 Related Work

Work in neural representations and language models has a long history (0; 0; 0) *inter alia*, and has recently regained a lot of attention.

(0) built a multi-task NLP system based on vector representation models which was trained to predict NER tags, chunkings, and semantic role labelings. The SRL task is the most relevant to this work, but their training procedure was fully supervised and is differs greatly from this work.

(0) showed that features learned by a few vector representation models improved the state of the art in NER tagging and chunking.

(0) analyzed the representations learned using a recurrent neural network language model described in (0) which was initially created to aid in machine translation. (0) however showed that the representations learned by the language model were not only very good for that task, but also for the task of analogy completion, which was one of the tasks in SemEval 2012 (0).

(0) presented a discriminatively trained feedforward neural network which induces representations with similar properties to those learned using a probabilistic language models, but with greatly accelerated training. This model achieved state of the art part of speech tagging accuracy (or near it) for many languages given only a very small labeled training set and the representations learned on unlabelled data. (0) also focussed on efficient training and proposed a similar model that sums rather than concatenates the input vectors, and another that is just the reverse prediction task with the same network structure (predict neighboring words from central word instead of central word from neighboring words). The representations worked very well at analogies that represent specific lexical shifts like pluralization (e.g. “dog” : “dogs” :: “cat” : “cats”), superlativization (e.g. “easy” : “easiest” :: “lucky” : “luckiest”), and capital cities (e.g. “Athens” : “Greece” :: “Oslo” : “Norway”). Their best model got 50.0% accuracy

on these tasks, but only one of the analogy types reflected a change in syntactic category, which was between adjectives and adverbs.

7 Conclusion

We present a basic model of coherence that induces features for words which shed light on their syntactic and conceptual properties.

References

- Rami Al-Rfou', Bryan Perozzi, and Steven Skiena. Polyglot: Distributed word representations for multilingual nlp. In *Proceedings of the Seventeenth Conference on Computational Natural Language Learning*, pages 183–192, Sofia, Bulgaria, August 2013. Association for Computational Linguistics.
- James Bergstra, Olivier Breuleux, Frédéric Bastien, Pascal Lamblin, Razvan Pascanu, Guillaume Desjardins, Joseph Turian, David Warde-Farley, and Yoshua Bengio. Theano: a CPU and GPU math expression compiler. In *Proceedings of the Python for Scientific Computing Conference (SciPy)*, June 2010. Oral Presentation.
- Ronan Collobert and Jason Weston. A unified architecture for natural language processing: deep neural networks with multitask learning. In William W. Cohen, Andrew McCallum, and Sam T. Roweis, editors, *ICML*, volume 307 of *ACM International Conference Proceeding Series*, pages 160–167. ACM, 2008.
- John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for online learning and stochastic optimization. *J. Mach. Learn. Res.*, 12:2121–2159, July 2011.
- K. Hornik, M. Stinchcombe, and H. White. Multilayer feedforward networks are universal approximators. *Neural Netw.*, 2(5):359–366, July 1989.
- Eric H. Huang, Richard Socher, Christopher D. Manning, and Andrew Y. Ng. Improving Word Representations via Global Context and Multiple Word Prototypes. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2012.
- David Jurgens, Saif Mohammad, Peter Turney, and Keith Holyoak. Semeval-2012 task 2: Measuring degrees of relational similarity. In **SEM 2012: The First Joint Conference on Lexical and Computational Semantics – Volume 1: Proceedings of the main conference and the shared task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation (SemEval 2012)*, pages 356–364, Montréal, Canada, 7–8 June 2012. Association for Computational Linguistics.
- Stefan Kombrink, Tomas Mikolov, Martin Karafiát, and Lukás Burget. Recurrent neural network based language modeling in meeting recognition. In *INTER-SPEECH*, pages 2877–2880. ISCA, 2011.
- Catherine Macleod, Ralph Grishman, Adam Meyers, Leslie Barrett, and Ruth Reeves. Nomlex: A lexicon of nominalizations. In *In Proceedings of Euralex98*, pages 187–193, 1998.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *CoRR*, abs/1301.3781, 2013.
- Tomas Mikolov, Wen tau Yih, and Geoffrey Zweig. Linguistic regularities in continuous space word representations. In *HLT-NAACL*, pages 746–751. The Association for Computational Linguistics, 2013.
- Courtney Napoles, Matthew Gormley, and Benjamin Van Durme. Annotated gigaword. In *AKBC-WEKEX Workshop at NAACL 2012*, June 2012.
- Jordan B. Pollack. Recursive distributed representations. *Artificial Intelligence*, 46:77–105, 1990.
- Joseph Reisinger and Raymond J. Mooney. Multi-prototype vector-space models of word meaning. In *Proceedings of the 11th Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-2010)*, pages 109–117, 2010.
- D. E. Rumelhart, G. E. Hinton, and R. J. Williams. Neurocomputing: Foundations of research. chapter Learning Internal Representations by Error Propagation, pages 673–695. MIT Press, Cambridge, MA, USA, 1988.
- David E. Rumelhart, Paul Smolensky, James L. McClelland, and Geoffrey E. Hinton. Schemata and sequential thought processes in PDP models. In James L. McClelland and David E. Rumelhart, editors, *Parallel Distributed Processing: Explorations in the Microstructure of Cognition, Volume 2: Psychological and Biological Models*, pages 7–57. MIT Press, Cambridge, MA, 1986.
- Joseph Turian, Lev Ratinov, and Yoshua Bengio. Word representations: A simple and general method for semi-supervised learning. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics, ACL '10*, pages 384–394, Stroudsburg, PA, USA, 2010. Association for Computational Linguistics.